

INTRODUCTION TO REGRESSION ANALYSIS

Bethany Baker

DS-SF-33//March 27, 2017

INTRODUCTION TO REGRESSION ANALYSIS

LEARNING OBJECTIVES

- Define data modeling and simple linear regression
- Build a linear regression model using a dataset that meets the linearity assumption using the sci-kit learn library
- Understand and identify multicollinearity in a multiple regression.

INTRODUCTION TO REGRESSION ANALYSIS

PRE-WORK

PRE-WORK REVIEW

- Recall in last class that we completed our review of Statistics

Fundamentals, consisting of:

- **Descriptive Statistics** (mode, median, mean, variance, etc.)
- **Inferential Statistics** (estimation and hypothesis testing) - reach generalizations about the whole by examining a sample

HYPOTHESIS TESTING

- *Testing Hypotheses* about population characteristics (such as the mean and standard deviation) is performed by first making an assumption about the unknown population characteristic
- We then take a random sample, and on the basis of the corresponding sample characteristic, either accept or reject the hypothesis with a particular degree of confidence

HYPOTHESIS TESTING

On to the White Board!

OPENING

INTRODUCTION TO REGRESSION ANALYSIS

BUILDING MODELS OF DATA


- Fundamentally, machine learning involved building mathematical models to help understand data
- “Learning” occurs when we give these models tunable parameters that can be adapted to observed data
- Once the models are “fit” to previously seen data, they can be used to predict and understand aspect of newly observed data

BUILDING MODELS OF DATA

Supervised Learning	Unsupervised Learning
<ul style="list-style-type: none">• Classification: labels are discrete categories• Regression: labels are continuous quantities	<p>“Letting the dataset speak for itself”</p> <ul style="list-style-type: none">• Clustering• Dimensionality Reduction
<ul style="list-style-type: none">• Linear Regression• Random Forests• Support Vector Machines	<ul style="list-style-type: none">• k-Means• PCA• Apriori

BUILDING MODELS OF DATA

STARTING WITH THE FUNDAMENTALS!



Supervised Learning	Unsupervised Learning
<ul style="list-style-type: none">• Classification: labels are discrete categories• Regression: labels are continuous quantities	<p>“Letting the dataset speak for itself”</p> <ul style="list-style-type: none">• Clustering• Dimensionality Reduction
<ul style="list-style-type: none">• Linear Regression• Random Forests• Support Vector Machines	<ul style="list-style-type: none">• k-Means• PCA• Apriori

WHY LINEAR REGRESSION?

- widely used
- runs fast
- easy to use (no tuning is required)
- highly interpretable
- basis for many other methods

WHERE ARE WE IN THE DATA SCIENCE WORKFLOW?


- Data has been **acquired** and **parsed**.
- Today we'll **refine** the data and **build** models.
- We'll also use plots to **represent** the results.

INTRODUCTION

SIMPLE LINEAR REGRESSION

SIMPLE LINEAR REGRESSION

- **Definition:** Explanation of a continuous variable given a series of independent variables
- The simplest version is just a line of best fit:
 $y = mx + b$
- Explain the relationship between **x** and **y** using the intercept **b** and the slope **m**.



SIMPLE LINEAR REGRESSION

- However, linear regression uses linear algebra to explain the relationship between *multiple* x's and y.
- The more sophisticated version: $y = \text{beta} * X + \text{alpha} (+ \text{error})$
- Explain the relationship between the matrix **X** and a dependent vector **y** using a y-intercept **alpha** and the relative coefficients **beta**.

SIMPLE LINEAR REGRESSION

- Linear regression works **best** when:
 - The data is normally distributed (but doesn't have to be)
 - X's significantly explain y (have low p-values)
 - X's are independent of each other (low multicollinearity)
 - Resulting values pass linear assumption (depends upon problem)
- If data is not normally distributed, we could introduce *bias*.

DEMO

REGRESSING AND NORMAL DISTRIBUTIONS

DEMO: REGRESSING AND NORMAL DISTRIBUTIONS

- Follow along with your starter code notebook while I walk through these examples.
- The first plot shows a relationship between two values, though not a linear solution.
- Note that `lmpplot()` returns a straight line plot.
- However, we can transform the data, both log-log distributions to get a linear solution.

GUIDED PRACTICE

USING SEABORN TO GENERATE SIMPLE LINEAR MODEL PLOTS

ACTIVITY: GENERATE SINGLE VARIABLE LINEAR MODEL PLOTS



EXERCISE

DIRECTIONS (15 minutes)

1. Update and complete the code in the starter notebook to use **lmplo** and display correlations between body weight and two dependent variables: **sleep_rem** and **awake**.

DELIVERABLE

Two plots

INTRODUCTION

SIMPLE REGRESSION ANALYSIS IN SKLEARN

SIMPLE LINEAR REGRESSION ANALYSIS IN SKLEARN

- Sklearn defines models as *objects* (in the OOP sense).
- You can use the following principles:
 - All sklearn modeling classes are based on the [base estimator](#). This means all models take a similar form.
 - All estimators take a matrix \mathbf{X} , either sparse or dense.
 - Supervised estimators also take a vector \mathbf{y} (the response).
 - Estimators can be customized through setting the appropriate parameters.

CLASSES AND OBJECTS IN OBJECT ORIENTED PROGRAMMING

- **Classes** are an abstraction for a complex set of ideas, e.g. *human*.
- Specific **instances** of classes can be created as **objects**.
 - *john_smith = human()*
- Objects have **properties**. These are attributes or other information.
 - *john_smith.age*
 - *john_smith.gender*
- Objects have **methods**. These are procedures associated with a class/object.
 - *john_smith.breathe()*
 - *john_smith.walk()*

SIMPLE LINEAR REGRESSION ANALYSIS IN SKLEARN

- General format for sklearn model classes and methods

```
# generate an instance of an estimator class
estimator = base_models.AnySKLearnObject()
# fit your data
estimator.fit(X, y)
# score it with the default scoring method (recommended to use the metrics module in the future)
estimator.score(X, y)
# predict a new set of data
estimator.predict(new_X)
# transform a new X if changes were made to the original X while fitting
estimator.transform(new_X)
```

- LinearRegression() doesn't have a transform function
- With this information, we can build a simple process for linear regression.

DEMO

SIGNIFICANCE IS KEY

DEMO: SIGNIFICANCE IS KEY

- Follow along with your starter code notebook while I walk through these examples.
- What does the residual plot tell us?
- How can we use the linear assumption?

GUIDED PRACTICE

USING THE LINEAR REGRESSION OBJECT

ACTIVITY: USING THE LINEAR REGRESSION OBJECT

DIRECTIONS (15 minutes)



EXERCISE

1. With a partner, generate two more models using the log-transformed data to see how this transform changes the model's performance.
2. Use the code on the following slide to complete #1.

DELIVERABLE

Two new models

ACTIVITY: USING THE LINEAR REGRESSION OBJECT



EXERCISE

DIRECTIONS (15 minutes)

```
X =  
y =  
loop = []  
for boolean in loop:  
    print 'y-intercept:', boolean  
    lm =  
linear_model.LinearRegression(fit_intercept=boolean)  
    get_linear_model_metrics(X, y, lm)  
    print
```

DELIVERABLE

Two new models

INDEPENDENT PRACTICE

BASE LINEAR REGRESSION CLASSES

ACTIVITY: BASE LINEAR REGRESSION CLASSES



EXERCISE

DIRECTIONS (20 minutes)

1. Experiment with the model evaluation function we have (**get_linear_model_metrics**) with the following sklearn estimator classes.
 - a. `linear_model.Lasso()`
 - b. `linear_model.Ridge()`
 - c. `linear_model.ElasticNet()`

Note: We'll cover these new regression techniques in a later class.

DELIVERABLE

New models and evaluation metrics

INTRODUCTION


MULTIPLE REGRESSION ANALYSIS

MULTIPLE REGRESSION ANALYSIS

- Simple linear regression with one variable can explain some variance, but using multiple variables can be much more powerful.
- We want our multiple variables to be mostly independent to avoid multicollinearity.
- Multicollinearity, when two or more variables in a regression are highly correlated, can cause problems with the model.

BIKE DATA EXAMPLE

- ▶ We can look at a correlation matrix of our bike data.
- ▶ Even if adding correlated variables to the model improves overall variance, it can introduce problems when explaining the output of your model.
- ▶ What happens if we use a second variable that isn't highly correlated with temperature?



GUIDED PRACTICE

MULTICOLLINEARITY WITH DUMMY VARIABLES

ACTIVITY: MULTICOLLINEARITY WITH DUMMY VARIABLES

DIRECTIONS (15 minutes)



EXERCISE

1. Load the bike data.
2. Run through the code on the following slide.
3. What happens to the coefficients when you include all weather situations instead of just including all except one?

DELIVERABLE

Two models' output

ACTIVITY: MULTICOLLINEARITY WITH DUMMY VARIABLES



EXERCISE

DIRECTIONS (15 minutes)

```
lm = linear_model.LinearRegression()
weather = pd.get_dummies(bike_data.weathersit)
get_linear_model_metrics(weather[[1, 2, 3, 4]], y, lm)
print

# drop the least significant, weather situation = 4
get_linear_model_metrics(weather[[1, 2, 3]], y, lm)
```

DELIVERABLE

Two models' output

GUIDED PRACTICE

COMBINING FEATURES INTO A BETTER MODEL

ACTIVITY: COMBINING FEATURES INTO A BETTER MODEL



EXERCISE

DIRECTIONS (15 minutes)

1. With a partner, complete the code on the following slide.
2. Visualize the correlations of all the numerical features built into the dataset.
3. Add the three significant weather situations into our current model.
4. Find two more features that are not correlated with the current features, but could be strong indicators for predicting guest riders.

DELIVERABLE

Visualization of correlations, new models

ACTIVITY: COMBINING FEATURES INTO A BETTER MODEL



EXERCISE

DIRECTIONS (15 minutes)

```
lm = linear_model.LinearRegression()
bikemodel_data = bike_data.join() # add in the three weather situations

cmap = sns.diverging_palette(220, 10, as_cmap=True)
correlations = # what are we getting the correlations of?
print correlations
print sns.heatmap(correlations, cmap=cmap)

columns_to_keep = [] #[which_variables?]
final_feature_set = bikemodel_data[columns_to_keep]

get_linear_model_metrics(final_feature_set, y, lm)
```

DELIVERABLE

Visualization of correlations, new models

INDEPENDENT PRACTICE

BUILDING MODELS FOR OTHER Y VARIABLES

ACTIVITY: BUILDING MODELS FOR OTHER Y VARIABLES



EXERCISE

DIRECTIONS (25 minutes)

1. Build a new model using a new y variable: registered riders.
2. Pay attention to the following:
 - a. the distribution of riders (should we rescale the data?)
 - b. checking correlations between the variables and y variable
 - c. choosing features to avoid multicollinearity
 - d. model complexity vs. explanation of variance
 - e. the linear assumption

BONUS

1. Which variables make sense to dummy?
2. What features might explain ridership but aren't included? Can you build these features with the included data and pandas?

DELIVERABLE

A new model and evaluation metrics

CONCLUSION

TOPIC REVIEW

CONCLUSION

- You should now be able to answer the following questions:
 - What is simple linear regression?
 - What makes multi-variable regressions more useful?
 - What challenges do they introduce?
 - How do you dummy a category variable?
 - How do you avoid a singular matrix?

WEEK 3 : LESSON 6

UPCOMING WORK

UPCOMING WORK

Week 4 : Lesson 8

- Project: Final Project, Deliverable 1
- Unit Project 2

INTRODUCTION TO REGRESSION ANALYSIS

Q & A

INTRODUCTION TO REGRESSION ANALYSIS

EXIT TICKET

DON'T FORGET TO FILL OUT YOUR EXIT TICKET!